

Cost is a cocycle

No Lean in this one — it is design theory ahead of the mechanization, and it wears the honest badges to prove it. Its claims are exactly the ones a future track should turn green.

The machine’s original cost model had a property so simple it was easy to mistake for essential: the cost of a run depended only on the **word** of instructions, not on the configuration it ran from. Cost was a homomorphism from the free monoid — cost of a concatenation is the sum, end of story. Two amendments now push against this: free up kept the property, but the proposed **dirty** pricing (a write-back is paid only when leaving a subtree that was written) does not: the same word can cost differently from different configurations. Is that the end of “algebraic”?

It is not, and this exposition says precisely what replaces it.

1 The cocycle law is functoriality

The machine’s semantics is a (partial) action of the word monoid M on the set of configurations. A monoid is a one-object category BM , and the action is a functor $BM \rightarrow \text{Set}$: the single object goes to the configuration set, each word to its (partial) transition function. The **action category** $\mathcal{C}$ is the Grothendieck construction — the category of elements — of this functor. Unwinding the definition: objects are configurations, and a morphism out of c is **not** a bare word but a pair (w, c) — a word runnable at c , welded to the configuration it runs from — with target $c \cdot w$ and composition

$$(w_2, c \cdot w_1) \circ (w_1, c) = (w_1 w_2, c).$$

The welding is the entire content of the construction: “the same word from a different configuration” becomes a **different morphism**, which is exactly the room a state-dependent cost needs to live in.

Now let $B\mathbb{N}$ be the one-object category with morphisms $(\mathbb{N}, +, 0)$. A cost model is a functor

$$\text{cost} : \mathcal{C} \rightarrow B\mathbb{N}.$$

Because a morphism of $\mathcal{C}$ is a pair, a cost model is forced to be a function of pairs; write $\text{cost}_c(w)$ for its value on the morphism (w, c) — subscript the source, nothing ad hoc. Functoriality over the composition law above unpacks to exactly two laws:

$$\text{cost}(\text{id}_c) = 0, \quad \text{cost}_c(w_1 w_2) = \text{cost}_c(w_1) + \text{cost}_{c \cdot w_1}(w_2).$$

The shifted subscript is not an extra axiom — it only records that the second arrow of the composite starts where the first one landed. The second law is the **cocycle law** — the name comes from group cohomology (for a group action the category of elements is the translation groupoid, and functions with this law are its 1-cocycles) and from ergodic theory, where “cocycles over a dynamical system” are classical — including, pleasingly, cocycles over Vershik’s adic transformations: the odometer keeps finding us.

The point of the definition: **state-dependent cost composes perfectly**. Nothing about the categorical semantics gets harder — a costed run is a morphism in the Writer-style graded setting (pairs of state-transform and price, composition adds prices), which is precisely the cost-graded monoid action the founding conversation already planned. The finite-presentation story also survives verbatim: define the cost table on finitely many cases — generator $\times$ **local observation** (the same locality read-branching already uses; dirtiness is one honest bit per node) — and functoriality extends it uniquely. The real cleanliness boundary was never “cost ignores state”; it is “cost reads only the local observation, from a finite table.”

Static prices are a factorization. The Grothendieck construction comes with a projection $\pi : \mathcal{C} \rightarrow BM$, $(w, c) \mapsto w$ — forget the state, keep the word. A cost model is word-only precisely when it factors as $\mathcal{C} \rightarrow BM \rightarrow B\mathbb{N}$ through π : functors on BM are monoid homomorphisms into $(\mathbb{N}, +)$, the static prices; functors on $\mathcal{C}$ are the state-dependent ones. “Cost ignores state” stops being a property you check equation by equation and becomes a shape — does the functor factor? — and the desk theorem below says that write-back cost fails to factor even after correction by a coboundary.

2 Coboundaries are potential functions

Inside the cocycles live two distinguished classes. A cocycle is a **homomorphism** when it factors through π — cost blind to state; the base machine and the mount-distance surcharge both live here. And a cocycle is a **coboundary** when it is $(\delta\Phi)_c(w) = \Phi(c \cdot w) - \Phi(c)$ for a potential Φ on configurations. Coboundaries are cohomology’s “trivial” cocycles — and they are exactly the potential functions of amortized analysis. Two cost models differing by a coboundary agree on every closed walk and differ only by boundary terms:

the cohomology class is the amortized cost.

Amortization-as-normalization (convo [21]) becomes: choose the cleanest representative of your class.

LEAN $\vdash$ Adic.Dyadic.freeUp_exchange

Free up is a change of representative

Let Φ = current depth of the head. On movement words, $2 \cdot \text{cost}_{\text{free}} + \Phi(\text{initial}) = \text{cost}_{\text{sym}} + \Phi(\text{final})$ exactly (down: $2 = 1 + 1$; up: $0 = 1 - 1$). The free-up amendment did not change the cohomology class of movement cost — it chose a cleaner representative, and the boundary term is the familiar “moves $\leq 2 \cdot \text{cost} + \text{initial depth}$ ” safety argument. Lean proves this by induction on words.

```

theorem freeUp_exchange {n : Nat} (word : Word) (head final : Head n)
  (hrun : run word head = some final) :
  2 * cost word + depth head = symCost word + depth final
axioms: Quot.sound, propext    hash: 0dc16bd7115b

```

3 Bandwidth is a nontrivial class (desk-proved)

Now the dirty model: down fills, a clean up discards for free, a dirty up pays its write-back. This cocycle reads one bit of local state, so it is **not** a homomorphism. The sharp question is whether it is secretly trivial — cohomologous to some word-only cost, its state-dependence absorbable into a potential.

DESK-PROVED — not mechanized

The write-back cocycle is not cohomologous to any homomorphism

No word-homomorphism h and potential Φ satisfy $\text{cost}_{\text{dirty}} = h + \delta\Phi$. *Desk argument (upgraded from an open claim, 2026-08-10; AIP-5 §2).* (i) $B\mathbb{N}$ is **commutative**, so any homomorphism factors through the abelianization: h sees only a word’s letter counts — how many down0s, down1s, ups, writes — nothing of their order. (ii) Build two **closed** walks at the same configuration c_0 (root, clean tree, all-zero store) with equal letter counts. The sparse walk makes m isolated round trips to depth- n leaves writing 1s, then repeats them writing 0s back — store restored, every path flushed. The dense walk writes a 2^k -leaf block ($2^k = m$) in one streaming sweep and erases it likewise, then **pads** with clean write-free excursions in an untouched subtree until every letter count matches — always solvable in root-returning chunks of depth $\leq n$, since each walk individually balances downs against ups. (iii) Identical endpoints kill $\delta\Phi$ **exactly** (both boundary terms are $\Phi(c_0) - \Phi(c_0) = 0$); equal counts make h agree; but the sparse walk pays $\Theta(mn)$ in write-backs (every up on an isolated dirty path pays, both passes) against the dense walk’s $\Theta(m + n)$ (each region edge dirty-crossed once; the padding pays its downs — which h sees equally on both sides — but its ups are clean and free). With $m = \Theta(n)$ the totals are $4mn + 2m$ against $2mn + 2m + \Theta(m + n)$: an unexplained $\Theta(n^2)$ gap — contradiction. The exact padding counts are written out in AIP-5 §2; the gap is precisely the sparse walk’s dirty ups against the padding’s clean ones.

So — modulo that bookkeeping — it is a theorem, not a slogan: **bandwidth cannot be statically priced**. Latency-like costs (addressing, movement, mount distance) are homomorphisms; write-back cost is genuinely history-dependent, and the cohomological language says so exactly. That real machines struggle to price bandwidth statically (rooflines, write amplification, “it depends on your access pattern”) stops being folklore and becomes the statement that a certain cohomology class is nonzero.

DESK-PROVED — not mechanized

Sparse–dense write asymmetry

Under dirty pricing: m isolated writes at pairwise tree-distant leaves cost $\Theta(m \cdot n)$ total (each pays its path home), while writing a full 2^k -subtree costs $\Theta(2^k)$ — each edge of the written region is dirty-crossed once, amortized $O(1)$ per written leaf. The write-side twin of the boundary spike; desk-level pending the dirty-bit mechanization.

4 What this buys the programme

Three earlier decisions become one picture. The machine’s semantics was already a (partial) monoid action; its cost was already meant to be a grading (convo [7], [13]); amortization was already normalization (convo [21]). The cocycle view identifies all three: **a cost model is a functor to $B\mathbb{N}$; static prices are the homomorphisms; potentials are the coboundaries; amortized equivalence is cohomology; and the interesting resources are the nontrivial classes**. The mechanization needs none of this vocabulary — it proves cost equations word by word, as it already does — but the paper gets to say why those equations have the shape they have.